Package 'DatabaseConnector'

May 9, 2023

Type Package

Title Connecting to Various Database Platforms

Version 6.2.1

Date 2023-05-09

Description An R 'DataBase Interface' ('DBI') compatible interface to various database plat-
forms ('PostgreSQL', 'Oracle', 'Microsoft SQL Server',
'Amazon Redshift', 'Microsoft Parallel Database Warehouse', 'IBM Netezza', 'Apache Im-
pala', 'Google BigQuery', 'Snowflake', 'Spark', and 'SQLite'). Also includes support for
fetching data as 'Andromeda' objects. Uses either 'Java Database Connectiv-
ity' ('JDBC') or other 'DBI' drivers to connect to databases.

SystemRequirements Java (>= 8)

Depends R (>= 4.0.0)

Imports rJava,
SqlRender (>= 1.15.0),
methods,
stringr,
readr,
rlang,
utils,
DBI (>= 1.0.0),
urltools,
bit64,
checkmate,
digest,
dbplyr (>= 2.2.0)

Suggests aws.s3,
R.utils,
withr,
testthat,
DBItest,
knitr,
rmarkdown,
RSQLite,
ssh,
Andromeda,
dplyr,
RPostgres,
odbc,

duckdb,
pool,
ParallelLogger

License Apache License

VignetteBuilder knitr

URL https://ohdsi.github.io/DatabaseConnector/, https:
//github.com/OHDSI/DatabaseConnector

BugReports https://github.com/OHDSI/DatabaseConnector/issues

Copyright See file COPYRIGHTS

RoxygenNote 7.2.3

Roxygen list(markdown = TRUE)

Encoding UTF-8

R **topics documented:**

---

assertTempEmulationSchemaSet

Assert the temp emulation schema is set

Description

Asserts the temp emulation schema is set for DBMSs requiring temp table emulation.

If you know your code uses temp tables, it is a good idea to call this function first, so it can throw an informative error if the user forgot to set the temp emulation schema.

Usage

```
assertTempEmulationSchemaSet(
  dbms,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema")
)
```

Arguments

dbms   The type of DBMS running on the server. See connect() or createConnectionDetails()
                    for valid values.

tempEmulationSchema
                    The temp emulation schema specified by the user.

Value

Does not return anything. Throws an error if the DBMS requires temp emulation but the temp emulation schema is not set.

---

computeDataHash   *Compute hash of data*

Description

Compute a hash of the data in the database schema. If the data changes, this should produce a different hash code. Specifically, the hash is based on the field names, field types, and table row counts.

Usage

```
computeDataHash(connection, databaseSchema, tables = NULL, progressBar = TRUE)
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
databaseSchema	The name of the database schema. See details for platform-specific details.
tables	(Optional) A list of tables to restrict to.
progressBar	When true, a progress bar is shown based on the number of tables in the database schema.

Details

The databaseSchema argument is interpreted differently according to the different platforms: SQL Server and PDW: The databaseSchema schema should specify both the database and the schema, e.g. 'my_database.dbo'. Impala: the databaseSchema should specify the database. Oracle: The databaseSchema should specify the Oracle 'user'. All other : The databaseSchema should specify the schema.

Value

A string representing the MD5 hash code.

---

connect   *connect*

Description

Creates a connection to a database server .There are four ways to call this function:

- connect(dbms, user, password, server, port, extraSettings, oracleDriver, pathToDriver)
- connect(connectionDetails)
- connect(dbms, connectionString, pathToDriver))
- connect(dbms, connectionString, user, password, pathToDriver)

DBMS parameter details::
Depending on the DBMS, the function arguments have slightly different interpretations:
Oracle:

- user. The user name used to access the server
- password. The password for that user
- server. This field contains the SID, or host and servicename, SID, or TNSName: 'sid', 'host/sid', 'host/service name', or 'tnsname'
- port. Specifies the port on the server (default = 1521)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "(PROTOCOL=tcps)")
- oracleDriver. The driver to be used. Choose between "thin" or "oci".
- pathToDriver. The path to the folder containing the Oracle JDBC driver JAR files.

Microsoft SQL Server:

- user. The user used to log in to the server. If the user is not specified, Windows Integrated Security will be used, which requires the SQL Server JDBC drivers to be installed (see details below).
- password. The password used to log on to the server
- server. This field contains the host name of the server
- port. Not used for SQL Server
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "encrypt=true; trustServerCertificate=false;")
- pathToDriver. The path to the folder containing the SQL Server JDBC driver JAR files.

Microsoft PDW:

- user. The user used to log in to the server. If the user is not specified, Windows Integrated Security will be used, which requires the SQL Server JDBC drivers to be installed (see details below).
- password. The password used to log on to the server
- server. This field contains the host name of the server
- port. Not used for SQL Server
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "encrypt=true; trustServerCertificate=false;")
- pathToDriver. The path to the folder containing the SQL Server JDBC driver JAR files.

PostgreSQL:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5432)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "ssl=true")
- pathToDriver. The path to the folder containing the PostgreSQL JDBC driver JAR files.

Redshift:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5439)
- 'extraSettings The configuration settings for the connection (i.e. SSL Settings such as "ssl=true&sslfactory=com.a
- pathToDriver. The path to the folder containing the RedShift JDBC driver JAR files.

Netezza:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5480)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "ssl=true")
- pathToDriver. The path to the folder containing the Netezza JDBC driver JAR file (nzjdbc.jar).

Impala:

- user. The user name used to access the server
- password. The password for that user
- server. The host name of the server
- port. Specifies the port on the server (default = 21050)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "SSLKeyStorePwd=*****")
- pathToDriver. The path to the folder containing the Impala JDBC driver JAR files.

SQLite:

- server. The path to the SQLIte file.

Spark / Databricks:

Currently both JDBC and ODBC connections are supported for Spark. Set the connectionString argument to use JDBC, otherwise ODBC is used:

- connectionString. The JDBC connection string (e.g. something like 'jdbc:databricks://my-org.cloud.databricks.com:443/default;transportMode=http;ssl=1;AuthMech=3;httpPath=/sql/1.0/warehouses/abcde
- user. The user name used to access the server.
- password. The password for that user.
- server. The host name of the server (when using ODBC), e.g. 'my-org.cloud.databricks.com')
- port. Specifies the port on the server (when using ODBC)
- extraSettings. Additional settings for the ODBC connection, for example extraSettings = list(HTTPPath = "/sql/1.0/warehouses/abcde12345", SSL = 1, ThriftTransport = 2, AuthMech = 3)

Snowflake:

- connectionString. The connection string (e.g. starting with 'jdbc:snowflake://host:port/?db=database').
- user. The user name used to access the server.
- password. The password for that user.

Windows authentication for SQL Server::

To be able to use Windows authentication for SQL Server (and PDW), you have to install the JDBC driver. Download the version 9.2.0 .zip from Microsoft and extract its contents to a folder. In the extracted folder you will find the file sqljdbc_9.2/enu/auth/x64/mssql-jdbc_auth-9.2.0.x64.dll (64-bits) or ssqljdbc_9.2/enu/auth/x86/mssql-jdbc_auth-9.2.0.x86.dll (32-bits), which needs to be moved to location on the system path, for example to c:/windows/system32. If you not have write access to any folder in the system path, you can also specify the path to the folder containing the dll by setting the environmental variable PATH_TO_AUTH_DLL, so for example Sys.setenv("PATH_TO_AUTH_DLL" = "c:/temp") Note that the environmental variable needs to be set before calling connect() for the first time.

Arguments

connectionDetails

> An object of class connectionDetails as created by the createConnectionDetails() function.

dbms   The type of DBMS running on the server. Valid values are

- "oracle" for Oracle
- "postgresql" for PostgreSQL
- "redshift" for Amazon Redshift
- "sql server" for Microsoft SQL Server
- "pdw" for Microsoft Parallel Data Warehouse (PDW)

- "netezza" for IBM Netezza
- "bigquery" for Google BigQuery
- "sqlite" for SQLite
- "sqlite extended" for SQLite with extended types (DATE and DATETIME)
- "spark" for Spark
- "snowflake" for Snowflake

user	The user name used to access the server.
password	The password for that user.
server	The name of the server.
port	(optional) The port on the server to connect to.
extraSettings	(optional) Additional configuration settings specific to the database provider to configure things as security for SSL. For connections using JDBC these will be appended to end of the connection string. For connections using DBI, these settings will additionally be used to call dbConnect().
oracleDriver	Specify which Oracle drive you want to use. Choose between "thin" or "oci".
connectionString	
	The JDBC connection string. If specified, the server, port, extraSettings, and oracleDriver fields are ignored. If user and password are not specified, they are assumed to already be included in the connection string.
pathToDriver	Path to a folder containing the JDBC driver JAR files. See downloadJdbcDrivers() for instructions on how to download the relevant drivers.

Details

This function creates a connection to a database.

Value

An object that extends DBIConnection in a database-specific manner. This object is used to direct commands to the database engine.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost/postgres",
  user = "root",
  password = "xxx"
)
conn <- connect(connectionDetails)
dbGetQuery(conn, "SELECT COUNT(*) FROM person")
disconnect(conn)

conn <- connect(dbms = "sql server", server = "RNDUSRDHIT06.jnj.com")
dbGetQuery(conn, "SELECT COUNT(*) FROM concept")
disconnect(conn)

conn <- connect(
  dbms = "oracle",
  server = "127.0.0.1/xe",
  user = "system",
```

```
  password = "xxx",
  pathToDriver = "c:/temp"
)
dbGetQuery(conn, "SELECT COUNT(*) FROM test_table")
disconnect(conn)

conn <- connect(
  dbms = "postgresql",
  connectionString = "jdbc:postgresql://127.0.0.1:5432/cmd_database"
)
dbGetQuery(conn, "SELECT COUNT(*) FROM person")
disconnect(conn)

## End(Not run)
```

---

createConnectionDetails

createConnectionDetails

Description

Creates a list containing all details needed to connect to a database. There are three ways to call this function:

- `createConnectionDetails(dbms, user, password, server, port, extraSettings, oracleDriver, pathToDriver)`
- `createConnectionDetails(dbms, connectionString, pathToDriver)`
- `createConnectionDetails(dbms, connectionString, user, password, pathToDriver)`

DBMS parameter details::

Depending on the DBMS, the function arguments have slightly different interpretations:

Oracle:

- `user`. The user name used to access the server
- `password`. The password for that user
- `server`. This field contains the SID, or host and servicename, SID, or TNSName: 'sid', 'host/sid', 'host/service name', or 'tnsname'
- `port`. Specifies the port on the server (default = 1521)
- `extraSettings`. The configuration settings for the connection (i.e. SSL Settings such as "(PROTOCOL=tcps)")
- `oracleDriver`. The driver to be used. Choose between "thin" or "oci".
- `pathToDriver`. The path to the folder containing the Oracle JDBC driver JAR files.

Microsoft SQL Server:

- `user`. The user used to log in to the server. If the user is not specified, Windows Integrated Security will be used, which requires the SQL Server JDBC drivers to be installed (see details below).
- `password`. The password used to log on to the server
- `server`. This field contains the host name of the server
- `port`. Not used for SQL Server

- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "encrypt=true; trustServerCertificate=false;")
- pathToDriver. The path to the folder containing the SQL Server JDBC driver JAR files.

Microsoft PDW:

- user. The user used to log in to the server. If the user is not specified, Windows Integrated Security will be used, which requires the SQL Server JDBC drivers to be installed (see details below).
- password. The password used to log on to the server
- server. This field contains the host name of the server
- port. Not used for SQL Server
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "encrypt=true; trustServerCertificate=false;")
- pathToDriver. The path to the folder containing the SQL Server JDBC driver JAR files.

PostgreSQL:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5432)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "ssl=true")
- pathToDriver. The path to the folder containing the PostgreSQL JDBC driver JAR files.

Redshift:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5439)
- 'extraSettings The configuration settings for the connection (i.e. SSL Settings such as "ssl=true&sslfactory=com.ar
- pathToDriver. The path to the folder containing the RedShift JDBC driver JAR files.

Netezza:

- user. The user used to log in to the server
- password. The password used to log on to the server
- server. This field contains the host name of the server and the database holding the relevant schemas: host/database
- port. Specifies the port on the server (default = 5480)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "ssl=true")
- pathToDriver. The path to the folder containing the Netezza JDBC driver JAR file (nzjdbc.jar).

Impala:

- user. The user name used to access the server
- password. The password for that user
- server. The host name of the server
- port. Specifies the port on the server (default = 21050)
- extraSettings. The configuration settings for the connection (i.e. SSL Settings such as "SSLKeyStorePwd=*****")

- pathToDriver. The path to the folder containing the Impala JDBC driver JAR files.

SQLite:

- server. The path to the SQLIte file.

Spark / Databricks:

Currently both JDBC and ODBC connections are supported for Spark. Set the connectionString argument to use JDBC, otherwise ODBC is used:

- connectionString. The JDBC connection string (e.g. something like 'jdbc:databricks://my-org.cloud.databricks.com:443/default;transportMode=http;ssl=1;AuthMech=3;httpPath=/sql/1.0/warehouses/abc
- user. The user name used to access the server.
- password. The password for that user.
- server. The host name of the server (when using ODBC), e.g. 'my-org.cloud.databricks.com')
- port. Specifies the port on the server (when using ODBC)
- extraSettings. Additional settings for the ODBC connection, for example extraSettings = list(HTTPPath = "/sql/1.0/warehouses/abcde12345", SSL = 1, ThriftTransport = 2, AuthMech = 3)

Snowflake:

- connectionString. The connection string (e.g. starting with 'jdbc:snowflake://host:port/?db=database').
- user. The user name used to access the server.
- password. The password for that user.

Windows authentication for SQL Server::

To be able to use Windows authentication for SQL Server (and PDW), you have to install the JDBC driver. Download the version 9.2.0 .zip from Microsoft and extract its contents to a folder. In the extracted folder you will find the file sqljdbc_9.2/enu/auth/x64/mssql-jdbc_auth-9.2.0.x64.dll (64-bits) or ssqljdbc_9.2/enu/auth/x86/mssql-jdbc_auth-9.2.0.x86.dll (32-bits), which needs to be moved to location on the system path, for example to c:/windows/system32. If you not have write access to any folder in the system path, you can also specify the path to the folder containing the dll by setting the environmental variable PATH_TO_AUTH_DLL, so for example Sys.setenv("PATH_TO_AUTH_DLL" = "c:/temp") Note that the environmental variable needs to be set before calling connect() for the first time.

Arguments

dbms   The type of DBMS running on the server. Valid values are

- "oracle" for Oracle
- "postgresql" for PostgreSQL
- "redshift" for Amazon Redshift
- "sql server" for Microsoft SQL Server
- "pdw" for Microsoft Parallel Data Warehouse (PDW)
- "netezza" for IBM Netezza
- "bigquery" for Google BigQuery
- "sqlite" for SQLite
- "sqlite extended" for SQLite with extended types (DATE and DATETIME)
- "spark" for Spark
- "snowflake" for Snowflake

user   The user name used to access the server.

password   The password for that user.

server	The name of the server.
port	(optional) The port on the server to connect to.
extraSettings	(optional) Additional configuration settings specific to the database provider to configure things as security for SSL. For connections using JDBC these will be appended to end of the connection string. For connections using DBI, these settings will additionally be used to call dbConnect().
oracleDriver	Specify which Oracle drive you want to use. Choose between ″thin″ or ″oci″.
connectionString	
	The JDBC connection string. If specified, the server, port, extraSettings, and oracleDriver fields are ignored. If user and password are not specified, they are assumed to already be included in the connection string.
pathToDriver	Path to a folder containing the JDBC driver JAR files. See downloadJdbcDrivers() for instructions on how to download the relevant drivers.

Details

This function creates a list containing all details needed to connect to a database. The list can then be used in the connect() function.

It is highly recommended to use a secure approach to storing credentials, so not to have your credentials in plain text in your R scripts. The examples demonstrate how to use the keyring package.

Value

A list with all the details needed to connect to a database.

Examples

```
## Not run:
# Needs to be done only once on a machine. Credentials will then be stored in
# the operating system's secure credential manager:
keyring::key_set_with_value("server", password = "localhost/postgres")
keyring::key_set_with_value("user", password = "root")
keyring::key_set_with_value("password", password = "secret")

# Create connection details using keyring. Note: the connection details will
# not store the credentials themselves, but the reference to get the credentials.
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = keyring::key_get("server"),
  user = keyring::key_get("user"),
  password = keyring::key_get("password"),
)
conn <- connect(connectionDetails)
dbGetQuery(conn, "SELECT COUNT(*) FROM person")
disconnect(conn)

## End(Not run)
```

---

createDbiConnectionDetails

Create DBI connection details

Description

For advanced users only. This function will allow `DatabaseConnector` to wrap any DBI driver. Using a driver that `DatabaseConnector` hasn't been tested with may give unpredictable performance. Use at your own risk. No support will be provided.

Usage

```
createDbiConnectionDetails(dbms, drv, ...)
```

Arguments

dbms   The type of DBMS running on the server. Valid values are

- "oracle" for Oracle
- "postgresql" for PostgreSQL
- "redshift" for Amazon Redshift
- "sql server" for Microsoft SQL Server
- "pdw" for Microsoft Parallel Data Warehouse (PDW)
- "netezza" for IBM Netezza
- "bigquery" for Google BigQuery
- "sqlite" for SQLite
- "sqlite extended" for SQLite with extended types (DATE and DATETIME)
- "spark" for Spark
- "snowflake" for Snowflake

drv   An object that inherits from DBIDriver, or an existing DBIConnection object (in order to clone an existing connection).

...   authentication arguments needed by the DBMS instance; these typically include user, password, host, port, dbname, etc. For details see the appropriate DBIDriver

Value

A list with all the details needed to connect to a database.

---

createZipFile   *Compress files and/or folders into a single zip file*

Description

Compress files and/or folders into a single zip file

Usage

```
createZipFile(zipFile, files, rootFolder = getwd(), compressionLevel = 9)
```

Arguments

`zipFile`	The path to the zip file to be created.
`files`	The files and/or folders to be included in the zip file. Folders will be included recursively.
`rootFolder`	The root folder. All files will be stored with relative paths relative to this folder.
`compressionLevel`	
	A number between 1 and 9. 9 compresses best, but it also takes the longest.

Details

Uses Java's compression library to create a zip file. It is similar to `utils::zip`, except that it does not require an external zip tool to be available on the system path.

`DatabaseConnectorDriver`
Create a DatabaseConnectorDriver object

Description

Create a DatabaseConnectorDriver object

Usage

```
DatabaseConnectorDriver()
```

`dateAdd`   *Add an interval to a date*

Description

This function is provided primarily to be used together with `dbplyr` when querying a database. It will also work in `dplyr` against data frames.

Usage

```
dateAdd(interval, number, date)
```

Arguments

`interval`	Unit for the interval. Can be "day", "week", "month", "year".
`number`	The number of units to add to the date.
`date`	The date to add to.

Value

A new date.

Examples

```
dateAdd("day", 10, as.Date("2000-01-01"))
```

---

dateDiff   *Compute difference between dates*

Description

This function is provided primarily to be used together with dbplyr when querying a database. It will also work in dplyr against data frames.

Usage

```
dateDiff(interval, date1, date2)
```

Arguments

interval	Unit for the interval. Can be "day", "week", "month", "year".
date1	The first date.
date2	The second date.

Value

The numeric value of the difference.

Examples

```
dateDiff("day", as.Date("2000-01-01"), as.Date("2000-03-01"))
```

---

dateFromParts   *Construct a date from parts*

Description

This function is provided primarily to be used together with dbplyr when querying a database. It will also work in dplyr against data frames.

Usage

```
dateFromParts(year, month, day)
```

Arguments

year	The calendar year.
month	The calendar month (1 = January).
day	The day of the month.

Value

The date.

Examples

```
dateFromParts(2000, 1, 5)
```

---

day   *Extract the day from a date*

Description

This function is provided primarily to be used together with dbplyr when querying a database. It will also work in dplyr against data frames.

Usage

```
day(date)
```

Arguments

date   The date.

Value

The day

Examples

```
day(as.Date("2000-02-01"))
```

---

dbConnect,DatabaseConnectorDriver-method
                    Create a connection to a DBMS

Description

Connect to a database. This function is synonymous with the [connect()](#) function. except a dummy driver needs to be specified

Usage

```
## S4 method for signature 'DatabaseConnectorDriver'
dbConnect(drv, ...)
```

Arguments

drv   The result of the [DatabaseConnectorDriver()](#) function

...   Other parameters. These are the same as expected by the [connect()](#) function.

Value

Returns a DatabaseConnectorConnection object that can be used with most of the other functions in this package.

Examples

```
## Not run:
conn <- dbConnect(DatabaseConnectorDriver(),
  dbms = "postgresql",
  server = "localhost/ohdsi",
  user = "joe",
  password = "secret"
)
querySql(conn, "SELECT * FROM cdm_synpuf.person;")
dbDisconnect(conn)

## End(Not run)
```

---

dbGetInfo,DatabaseConnectorDriver-method
Get DBMS metadata

Description

Retrieves information on objects of class DBIDriver, DBIConnection or DBIResult.

Usage

```
## S4 method for signature 'DatabaseConnectorDriver'
dbGetInfo(dbObj, ...)
```

Arguments

dbObj   An object inheriting from DBIObject, i.e. DBIDriver, DBIConnection, or a DBIResult

...   Other arguments to methods.

Value

For objects of class DBIDriver, dbGetInfo() returns a named list that contains at least the following components:

- driver.version: the package version of the DBI backend,
- client.version: the version of the DBMS client library.

For objects of class DBIConnection, dbGetInfo() returns a named list that contains at least the following components:

- db.version: version of the database server,
- dbname: database name,
- username: username to connect to the database,

- host: hostname of the database server,
- port: port on the database server. It must not contain a password component. Components that are not applicable should be set to NA.

For objects of class [DBIResult](#), dbGetInfo() returns a named list that contains at least the following components:

- statatment: the statement used with [dbSendQuery()](#) or [dbExecute()](#), as returned by [dbGetStatement()](#),
- row.count: the number of rows fetched so far (for queries), as returned by [dbGetRowCount()](#),
- rows.affected: the number of rows affected (for statements), as returned by [dbGetRowsAffected()](#)
- has.completed: a logical that indicates if the query or statement has completed, as returned by [dbHasCompleted()](#).

See Also

Other DBIDriver generics: [DBIDriver-class](#), [dbCanConnect()](#), [dbConnect()](#), [dbDataType()](#), [dbDriver()](#), [dbIsReadOnly()](#), [dbIsValid()](#), [dbListConnections()](#)

Other DBIConnection generics: [DBIConnection-class](#), [dbAppendTable()](#), [dbCreateTable()](#), [dbDataType()](#), [dbDisconnect()](#), [dbExecute()](#), [dbExistsTable()](#), [dbGetException()](#), [dbGetQuery()](#), [dbIsReadOnly()](#), [dbIsValid()](#), [dbListFields()](#), [dbListObjects()](#), [dbListResults()](#), [dbListTables()](#), [dbReadTable()](#), [dbRemoveTable()](#), [dbSendQuery()](#), [dbSendStatement()](#), [dbWriteTable()](#)

Other DBIResult generics: [DBIResult-class](#), [dbBind()](#), [dbClearResult()](#), [dbColumnInfo()](#), [dbFetch()](#), [dbGetRowCount()](#), [dbGetRowsAffected()](#), [dbGetStatement()](#), [dbHasCompleted()](#), [dbIsReadOnly()](#), [dbIsValid()](#), [dbQuoteIdentifier()](#), [dbQuoteLiteral()](#), [dbQuoteString()](#), [dbUnquoteIdentifier()](#)

---

dbms *Get the database platform from a connection*

Description

The SqlRender package provides functions that translate SQL from OHDSI-SQL to a target SQL dialect. These function need the name of the database platform to translate to. The dbms function returns the dbms for any DBI connection that can be passed along to SqlRender translation functions (see example).

Usage

```
dbms(connection)
```

Arguments

connection   The connection to the database server created using either [connect()](#) or [dbConnect()](#).

Value

The name of the database (dbms) used by SqlRender

Examples

```
library(DatabaseConnector)
con <- connect(dbms = "sqlite", server = ":memory:")
dbms(con)
#> [1] "sqlite"
SqlRender::translate("DATEADD(d, 365, dateColumn)", targetDialect = dbms(con))
#> "CAST(STRFTIME('%s', DATETIME(dateColumn, 'unixepoch', (365)||' days')) AS REAL)"
disconnect(con)
```

---

dbUnloadDriver,DatabaseConnectorDriver-method
Load and unload database drivers

Description

These methods are deprecated, please consult the documentation of the individual backends for the construction of driver instances.

dbDriver() is a helper method used to create an new driver object given the name of a database or the corresponding R package. It works through convention: all DBI-extending packages should provide an exported object with the same name as the package. dbDriver() just looks for this object in the right places: if you know what database you are connecting to, you should call the function directly.

dbUnloadDriver() is not implemented for modern backends.

Usage

```
## S4 method for signature 'DatabaseConnectorDriver'
dbUnloadDriver(drv, ...)
```

Arguments

drv	an object that inherits from DBIDriver as created by dbDriver.
...	any other arguments are passed to the driver drvName.

Details

The client part of the database communication is initialized (typically dynamically loading C code, etc.) but note that connecting to the database engine itself needs to be done through calls to dbConnect.

Value

In the case of dbDriver, an driver object whose class extends DBIDriver. This object may be used to create connections to the actual DBMS engine.

In the case of dbUnloadDriver, a logical indicating whether the operation succeeded or not.

See Also

Other DBIDriver generics: DBIDriver-class, dbCanConnect(), dbConnect(), dbDataType(), dbGetInfo(), dbIsReadOnly(), dbIsValid(), dbListConnections()

Other DBIDriver generics: DBIDriver-class, dbCanConnect(), dbConnect(), dbDataType(), dbGetInfo(), dbIsReadOnly(), dbIsValid(), dbListConnections()

---

disconnect   *Disconnect from the server*

Description

Close the connection to the server.

Usage

```
disconnect(connection)
```

Arguments

connection   The connection to the database server created using either [connect()](#) or [dbConnect()](#).

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah"
)
conn <- connect(connectionDetails)
count <- querySql(conn, "SELECT COUNT(*) FROM person")
disconnect(conn)

## End(Not run)
```

---

downloadJdbcDrivers   *Download DatabaseConnector JDBC Jar files*

Description

Download the DatabaseConnector JDBC drivers from https://ohdsi.github.io/DatabaseConnectorJars/

Usage

```
downloadJdbcDrivers(
  dbms,
  pathToDriver = Sys.getenv("DATABASECONNECTOR_JAR_FOLDER"),
  method = "auto",
  ...
)
```

Arguments

dbms	The type of DBMS to download Jar files for.

- "postgresql" for PostgreSQL
- "redshift" for Amazon Redshift
- "sql server", "pdw" or "synapse" for Microsoft SQL Server
- "oracle" for Oracle
- "spark" for Spark
- "snowflake" for Snowflake
- "bigquery" for Google BigQuery
- "all" for all aforementioned platforms

pathToDriver	The full path to the folder where the JDBC driver .jar files should be downloaded to. By default the value of the environment variable "DATABASECONNECTOR_JAR_FOLDER" is used.
method	The method used for downloading files. See ?download.file for details and options.
...	Further arguments passed on to download.file.

Details

The following versions of the JDBC drivers are currently used:

- PostgreSQL: V42.2.18
- RedShift: V2.1.0.9
- SQL Server: V9.2.0
- Oracle: V19.8
- Spark: V2.6.21
- Snowflake: V3.13.22
- BigQuery: v1.3.2.1003

Value

Invisibly returns the destination if the download was successful.

Examples

```
## Not run:
downloadJdbcDrivers("redshift")

## End(Not run)
```

dropEmulatedTempTables

Drop all emulated temp tables.

Description

On some DBMSs, like Oracle and BigQuery, `DatabaseConnector` through `SqlRender` emulates temp tables in a schema provided by the user. Ideally, these tables are deleted by the application / R script creating them, but for various reasons orphan temp tables may remain. This function drops all emulated temp tables created in this session only.

Usage

```
dropEmulatedTempTables(
  connection,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema")
)
```

Arguments

connection   The connection to the database server created using either [connect()](#) or [dbConnect()](#).

tempEmulationSchema

Some database platforms like Oracle and Impala do not truly support temp tables. To emulate temp tables, provide a schema with write privileges where temp tables can be created.

Value

Invisibly returns the list of deleted emulated temp tables.

eoMonth   *Return the end of the month*

Description

This function is provided primarily to be used together with `dbplyr` when querying a database. It will also work in `dplyr` against data frames.

Usage

```
eoMonth(date)
```

Arguments

date   A date in the month for which we need the end.

Value

The date of the last day of the month.

Examples

```
eoMonth(as.Date("2000-02-01"))
```

---

executeSql   *Execute SQL code*

Description

This function executes SQL consisting of one or more statements.

Usage

```
executeSql(
  connection,
  sql,
  profile = FALSE,
  progressBar = !as.logical(Sys.getenv("TESTTHAT", unset = FALSE)),
  reportOverallTime = TRUE,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  runAsBatch = FALSE
)
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
sql	The SQL to be executed
profile	When true, each separate statement is written to file prior to sending to the server, and the time taken to execute a statement is displayed.
progressBar	When true, a progress bar is shown based on the statements in the SQL code.
reportOverallTime	
	When true, the function will display the overall time taken to execute all statements.
errorReportFile	
	The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.
runAsBatch	When true the SQL statements are sent to the server as a single batch, and executed there. This will be faster if you have many small SQL statements, but there will be no progress bar, and no per-statement error messages. If the database platform does not support batched updates the query is executed without batching.

Details

This function splits the SQL in separate statements and sends it to the server for execution. If an error occurs during SQL execution, this error is written to a file to facilitate debugging. Optionally, a progress bar is shown and the total time taken to execute the SQL is displayed. Optionally, each separate SQL statement is written to file, and the execution time per statement is shown to aid in detecting performance issues.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
executeSql(conn, "CREATE TABLE x (k INT); CREATE TABLE y (k INT);")
disconnect(conn)

## End(Not run)
```

---

existsTable	*Does the table exist?*

Description

Checks whether a table exists. Accounts for surrounding escape characters. Case insensitive.

Usage

```
existsTable(connection, databaseSchema, tableName)
```

Arguments

connection	The connection to the database server created using either connect() or dbConnect().
databaseSchema	The name of the database schema. See details for platform-specific details.
tableName	The name of the table to check.

Details

The databaseSchema argument is interpreted differently according to the different platforms: SQL Server and PDW: The databaseSchema schema should specify both the database and the schema, e.g. 'my_database.dbo'. Impala: the databaseSchema should specify the database. Oracle: The databaseSchema should specify the Oracle 'user'. All other : The databaseSchema should specify the schema.

Value

A logical value indicating whether the table exits.

---

extractQueryTimes   *Extract query times from a* ParallelLogger *log file*

Description

When using the ParallelLogger default file logger, and using options(LOG_DATABASECONNECTOR_SQL = TRUE), DatabaseConnector will log all SQL sent to the server, and the time to get a response.

This function parses the log file, producing a data frame with time per query.

Usage

```
extractQueryTimes(logFileName)
```

Arguments

logFileName   Name of the ParallelLogger log file. Assumes the file was created using the default file logger.

Value

A data frame with queries and their run times in milliseconds.

Examples

```
connection <- connect(dbms = "sqlite", server = ":memory:")
logFile <- tempfile(fileext = ".log")
ParallelLogger::addDefaultFileLogger(fileName = logFile, name = "MY_LOGGER")
options(LOG_DATABASECONNECTOR_SQL = TRUE)

executeSql(connection, "CREATE TABLE test (x INT);")
querySql(connection, "SELECT * FROM test;")

extractQueryTimes(logFile)

ParallelLogger::unregisterLogger("MY_LOGGER")
unlink(logFile)
disconnect(connection)
```

---

getAvailableJavaHeapSpace
                    Get available Java heap space

Description

For debugging purposes: get the available Java heap space.

Usage

```
getAvailableJavaHeapSpace()
```

Value

The Java heap space (in bytes).

---

getTableNames *List all tables in a database schema.*

Description

This function returns a list of all tables in a database schema.

Usage

```
getTableNames(connection, databaseSchema = NULL, cast = "lower")
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
databaseSchema	The name of the database schema. See details for platform-specific details.
cast	Should the table names be cast to uppercase or lowercase before being returned? Valid options are "upper" , "lower" (default), "none" (no casting is done)

Details

The databaseSchema argument is interpreted differently according to the different platforms: SQL Server and PDW: The databaseSchema schema should specify both the database and the schema, e.g. 'my_database.dbo'. Impala: the databaseSchema should specify the database. Oracle: The databaseSchema should specify the Oracle 'user'. All other : The databaseSchema should specify the schema.

Value

A character vector of table names.

---

inDatabaseSchema *Refer to a table in a database schema*

Description

Can be used with [dplyr::tbl()](#) to indicate a table in a specific database schema.

Usage

```
inDatabaseSchema(databaseSchema, table)
```

Arguments

databaseSchema	The name of the database schema. See details for platform-specific details.
table	The name of the table in the database schema.

Details

The databaseSchema argument is interpreted differently according to the different platforms: SQL
Server and PDW: The databaseSchema schema should specify both the database and the schema,
e.g. 'my_database.dbo'. Impala: the databaseSchema should specify the database. Oracle: The
databaseSchema should specify the Oracle 'user'. All other : The databaseSchema should specify
the schema.

Value

An object representing the table and database schema.

---

insertTable   *Insert a table on the server*

Description

This function sends the data in a data frame to a table on the server. Either a new table is created,
or the data is appended to an existing table.

Usage

```
insertTable(
  connection,
  databaseSchema = NULL,
  tableName,
  data,
  dropTableIfExists = TRUE,
  createTable = TRUE,
  tempTable = FALSE,
  oracleTempSchema = NULL,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema"),
  bulkLoad = Sys.getenv("DATABASE_CONNECTOR_BULK_UPLOAD"),
  useMppBulkLoad = Sys.getenv("USE_MPP_BULK_LOAD"),
  progressBar = FALSE,
  camelCaseToSnakeCase = FALSE
)
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
databaseSchema	The name of the database schema. See details for platform-specific details.
tableName	The name of the table where the data should be inserted.
data	The data frame containing the data to be inserted.
dropTableIfExists	
	Drop the table if the table already exists before writing?
createTable	Create a new table? If false, will append to existing table.
tempTable	Should the table created as a temp table?
oracleTempSchema	
	DEPRECATED: use tempEmulationSchema instead.

`tempEmulationSchema`

Some database platforms like Oracle and Impala do not truly support temp tables. To emulate temp tables, provide a schema with write privileges where temp tables can be created.

`bulkLoad`   If using Redshift, PDW, Hive or Postgres, use more performant bulk loading techniques. Does not work for temp tables (except for HIVE). See Details for requirements for the various platforms.

`useMppBulkLoad`   DEPRECATED. Use `bulkLoad` instead.

`progressBar`   Show a progress bar when uploading?

`camelCaseToSnakeCase`

If TRUE, the data frame column names are assumed to use camelCase and are converted to snake_case before uploading.

Details

The `databaseSchema` argument is interpreted differently according to the different platforms: SQL Server and PDW: The databaseSchema schema should specify both the database and the schema, e.g. 'my_database.dbo'. Impala: the databaseSchema should specify the database. Oracle: The databaseSchema should specify the Oracle 'user'. All other : The databaseSchema should specify the schema.

This function sends the data in a data frame to a table on the server. Either a new table is created, or the data is appended to an existing table. NA values are inserted as null values in the database.

Bulk uploading:

Redshift: The MPP bulk loading relies upon the CloudyR S3 library to test a connection to an S3 bucket using AWS S3 credentials. Credentials are configured directly into the System Environment using the following keys: Sys.setenv("AWS_ACCESS_KEY_ID" = "some_access_key_id", "AWS_SECRET_ACCESS_KEY" = "some_secret_access_key", "AWS_DEFAULT_REGION" = "some_aws_region", "AWS_BUCKET_NAME" = "some_bucket_name", "AWS_OBJECT_KEY" = "some_object_key", "AWS_SSE_TYPE" = "server_side_encryption_type").

PDW: The MPP bulk loading relies upon the client having a Windows OS and the DWLoader exe installed, and the following permissions granted: –Grant BULK Load permissions - needed at a server level USE master; GRANT ADMINISTER BULK OPERATIONS TO user; –Grant Staging database permissions - we will use the user db. USE scratch; EXEC sp_addrolemember 'db_ddladmin', user; Set the R environment variable DWLOADER_PATH to the location of the binary.

PostgreSQL: Uses the 'psql' executable to upload. Set the POSTGRES_PATH environment variable to the Postgres binary path, e.g. 'C:/Program Files/PostgreSQL/11/bin'.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "mysql",
  server = "localhost",
  user = "root",
  password = "blah"
)
conn <- connect(connectionDetails)
data <- data.frame(x = c(1, 2, 3), y = c("a", "b", "c"))
insertTable(conn, "my_schema", "my_table", data)
disconnect(conn)
```

```
## bulk data insert with Redshift or PDW
connectionDetails <- createConnectionDetails(
  dbms = "redshift",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v5"
)
conn <- connect(connectionDetails)
data <- data.frame(x = c(1, 2, 3), y = c("a", "b", "c"))
insertTable(
  connection = connection,
  databaseSchema = "scratch",
  tableName = "somedata",
  data = data,
  dropTableIfExists = TRUE,
  createTable = TRUE,
  tempTable = FALSE,
  bulkLoad = TRUE
) # or, Sys.setenv("DATABASE_CONNECTOR_BULK_UPLOAD" = TRUE)

## End(Not run)
```

---

isSqlReservedWord   *Test a character vector of SQL names for SQL reserved words*

Description

This function checks a character vector against a predefined list of reserved SQL words.

Usage

```
isSqlReservedWord(sqlNames, warn = FALSE)
```

Arguments

sqlNames   A character vector containing table or field names to check.

warn   (logical) Should a warn be thrown if invalid SQL names are found?

Value

A logical vector with length equal to sqlNames that is TRUE for each name that is reserved and FALSE otherwise

---

jdbcDrivers	*How to download and use JDBC drivers for the various data platforms.*

Description

Below are instructions for downloading JDBC drivers for the various data platforms. Once downloaded use the pathToDriver argument in the connect() or createConnectionDetails() functions to point to the driver. Alternatively, you can set the 'DATABASECONNECTOR_JAR_FOLDER' environmental variable, for example in your .Renviron file (recommended).

SQL Server, Oracle, PostgreSQL, PDW, Snowflake, Spark, RedShift, Azure Synapse, BigQuery

Use the downloadJdbcDrivers() function to download these drivers from the OHDSI GitHub pages.

Netezza

Read the instructions here on how to obtain the Netezza JDBC driver.

Impala

Go to Cloudera's site, pick your OS version, and click "GET IT NOW!'. Register, and you should be able to download the driver.

SQLite

For SQLite we actually don't use a JDBC driver. Instead, we use the RSQLite package, which can be installed using install.packages("RSQLite").

---

lowLevelExecuteSql	*Execute SQL code*

Description

This function executes a single SQL statement.

Usage

```
lowLevelExecuteSql(connection, sql)
```

Arguments

connection	The connection to the database server created using either connect() or dbConnect().
sql	The SQL to be executed

lowLevelQuerySql   *Low level function for retrieving data to a data frame*

Description

This is the equivalent of the querySql() function, except no error report is written when an error occurs.

Usage

```
lowLevelQuerySql(
  connection,
  query,
  datesAsString = FALSE,
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE)
)
```

Arguments

connection	The connection to the database server created using either connect() or dbConnect().
query	The SQL statement to retrieve the data
datesAsString	Logical: Should dates be imported as character vectors, our should they be converted to R's date format?
integerAsNumeric	
	Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.
integer64AsNumeric	
	Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.

Details

Retrieves data from the database server and stores it in a data frame. Null values in the database are converted to NA values in R.

Value

A data frame containing the data retrieved from the server

lowLevelQuerySqlToAndromeda
                    Low level function for retrieving data to a local Andromeda object

Description

This is the equivalent of the querySqlToAndromeda() function, except no error report is written when an error occurs.

Usage

```
lowLevelQuerySqlToAndromeda(
  connection,
  query,
  andromeda,
  andromedaTableName,
  datesAsString = FALSE,
  appendToTable = FALSE,
  snakeCaseToCamelCase = FALSE,
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE)
)
```

Arguments

connection   The connection to the database server created using either [connect()](#) or [dbConnect()](#).

query   The SQL statement to retrieve the data

andromeda   An open Andromeda object, for example as created using [Andromeda::andromeda()](#).

andromedaTableName

        The name of the table in the local Andromeda object where the results of the query will be stored.

datesAsString   Should dates be imported as character vectors, our should they be converted to R's date format?

appendToTable   If FALSE, any existing table in the Andromeda with the same name will be replaced with the new data. If TRUE, data will be appended to an existing table, assuming it has the exact same structure.

snakeCaseToCamelCase

        If true, field names are assumed to use snake_case, and are converted to camelCase.

integerAsNumeric

        Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.

integer64AsNumeric

        Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.

Details

Retrieves data from the database server and stores it in a local Andromeda object This allows very large data sets to be retrieved without running out of memory. Null values in the database are converted to NA values in R. If a table with the same name already exists in the local Andromeda object it is replaced.

Value

Invisibly returns the andromeda. The Andromeda object will have a table added with the query results.

---

month   *Extract the month from a date*

Description

This function is provided primarily to be used together with dbplyr when querying a database. It will also work in dplyr against data frames.

Usage

```
month(date)
```

Arguments

date   The date.

Value

The month

Examples

```
month(as.Date("2000-02-01"))
```

---

querySql   *Retrieve data to a data.frame*

Description

This function sends SQL to the server, and returns the results.

Usage

```
querySql(
  connection,
  sql,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  snakeCaseToCamelCase = FALSE,
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE)
)
```

Arguments

connection	The connection to the database server created using either connect() or dbConnect().
sql	The SQL to be send.
errorReportFile	
	The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.
snakeCaseToCamelCase	
	If true, field names are assumed to use snake_case, and are converted to camelCase.
integerAsNumeric	
	Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.
integer64AsNumeric	
	Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.

Details

This function sends the SQL to the server and retrieves the results. If an error occurs during SQL execution, this error is written to a file to facilitate debugging. Null values in the database are converted to NA values in R.

Value

A data frame.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
count <- querySql(conn, "SELECT COUNT(*) FROM person")
disconnect(conn)

## End(Not run)
```

---

querySqlToAndromeda   *Retrieves data to a local Andromeda object*

Description

This function sends SQL to the server, and returns the results in a local Andromeda object

Usage

```
querySqlToAndromeda(
  connection,
  sql,
  andromeda,
  andromedaTableName,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  snakeCaseToCamelCase = FALSE,
  appendToTable = FALSE,
 integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE)
)
```

Arguments

connection   The connection to the database server created using either [connect()](#) or [dbConnect()](#).

sql   The SQL to be sent.

andromeda   An open Andromeda object, for example as created using [Andromeda::andromeda()](#).

andromedaTableName

The name of the table in the local Andromeda object where the results of the query will be stored.

errorReportFile

The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.

snakeCaseToCamelCase

If true, field names are assumed to use snake_case, and are converted to camel-Case.

appendToTable   If FALSE, any existing table in the Andromeda with the same name will be replaced with the new data. If TRUE, data will be appended to an existing table, assuming it has the exact same structure.

integerAsNumeric

Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.

integer64AsNumeric

Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.

Details

Retrieves data from the database server and stores it in a local Andromeda object. This allows very large data sets to be retrieved without running out of memory. If an error occurs during SQL execution, this error is written to a file to facilitate debugging. Null values in the database are converted to NA values in R.If a table with the same name already exists in the local Andromeda object it is replaced.

Value

Invisibly returns the andromeda. The Andromeda object will have a table added with the query results.

Examples

```
## Not run:
andromeda <- Andromeda::andromeda()
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
querySqlToAndromeda(
  connection = conn,
  sql = "SELECT * FROM person;",
  andromeda = andromeda,
  andromedaTableName = "foo"
)
disconnect(conn)

andromeda$foo

## End(Not run)
```

---

renderTranslateExecuteSql

Render, translate, execute SQL code

Description

This function renders, translates, and executes SQL consisting of one or more statements.

Usage

```
renderTranslateExecuteSql(
  connection,
  sql,
  profile = FALSE,
  progressBar = TRUE,
  reportOverallTime = TRUE,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  runAsBatch = FALSE,
  oracleTempSchema = NULL,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema"),
  ...
)
```

Arguments

connection	The connection to the database server created using either connect() or dbConnect().
sql	The SQL to be executed
profile	When true, each separate statement is written to file prior to sending to the server, and the time taken to execute a statement is displayed.

progressBar   When true, a progress bar is shown based on the statements in the SQL code.

reportOverallTime

                    When true, the function will display the overall time taken to execute all state-
                    ments.

errorReportFile

                    The file where an error report will be written if an error occurs. Defaults to
                    'errorReportSql.txt' in the current working directory.

runAsBatch   When true the SQL statements are sent to the server as a single batch, and exe-
                    cuted there. This will be faster if you have many small SQL statements, but there
                    will be no progress bar, and no per-statement error messages. If the database
                    platform does not support batched updates the query is executed as ordinarily.

oracleTempSchema

                    DEPRECATED: use `tempEmulationSchema` instead.

tempEmulationSchema

                    Some database platforms like Oracle and Impala do not truly support temp ta-
                    bles. To emulate temp tables, provide a schema with write privileges where
                    temp tables can be created.

...   Parameters that will be used to render the SQL.

Details

This function calls the `render` and `translate` functions in the `SqlRender` package before calling
`executeSql()`.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
renderTranslateExecuteSql(connection,
  sql = "SELECT * INTO #temp FROM @schema.person;",
  schema = "cdm_synpuf"
)
disconnect(conn)

## End(Not run)
```

---

renderTranslateQueryApplyBatched

                    Render, translate, and perform process to batches of data.

Description

This function renders, and translates SQL, sends it to the server, processes the data in batches with
a call back function. Note that this function should perform a row-wise operation. This is designed
to work with massive data that won't fit in to memory.

The batch sizes are determined by the java virtual machine and will depend on the data.

Usage

```
renderTranslateQueryApplyBatched(
  connection,
  sql,
  fun,
  args = list(),
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  snakeCaseToCamelCase = FALSE,
  oracleTempSchema = NULL,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema"),
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE),
  ...
)
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
sql	The SQL to be send.
fun	Function to apply to batch. Must take data.frame and integer position as parameters.
args	List of arguments to be passed to function call.
errorReportFile	
	The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.
snakeCaseToCamelCase	
	If true, field names are assumed to use snake_case, and are converted to camelCase.
oracleTempSchema	
	DEPRECATED: use `tempEmulationSchema` instead.
tempEmulationSchema	
	Some database platforms like Oracle and Impala do not truly support temp tables. To emulate temp tables, provide a schema with write privileges where temp tables can be created.
integerAsNumeric	
	Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.
integer64AsNumeric	
	Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.
...	Parameters that will be used to render the SQL.

Details

This function calls the `render` and `translate` functions in the `SqlRender` package before calling [querySql()](#).

Value

Invisibly returns a list of outputs from each call to the provided function.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
connection <- connect(connectionDetails)

# First example: write data to a large CSV file:
filepath <- "myBigFile.csv"
writeBatchesToCsv <- function(data, position, ...) {
  write.csv(data, filepath, append = position != 1)
  return(NULL)
}
renderTranslateQueryApplyBatched(connection,
  "SELECT * FROM @schema.person;",
  schema = "cdm_synpuf",
  fun = writeBatchesToCsv
)

# Second example: write data to Andromeda
# (Alternative to querySqlToAndromeda if some local computation needs to be applied)
bigResults <- Andromeda::andromeda()
writeBatchesToAndromeda <- function(data, position, ...) {
  data$p <- EmpiricalCalibration::computeTraditionalP(data$logRr, data$logSeRr)
  if (position == 1) {
    bigResults$rrs <- data
  } else {
    Andromeda::appendToTable(bigResults$rrs, data)
  }
  return(NULL)
}
sql <- "SELECT target_id, comparator_id, log_rr, log_se_rr FROM @schema.my_results;"
renderTranslateQueryApplyBatched(connection,
  sql,
  fun = writeBatchesToAndromeda,
  schema = "my_results",
  snakeCaseToCamelCase = TRUE
)

disconnect(connection)

## End(Not run)
```

---

renderTranslateQuerySql

Render, translate, and query to data.frame

Description

This function renders, and translates SQL, sends it to the server, and returns the results as a data.frame.

Usage

```
renderTranslateQuerySql(
  connection,
  sql,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  snakeCaseToCamelCase = FALSE,
  oracleTempSchema = NULL,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema"),
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE),
  ...
)
```

Arguments

connection	The connection to the database server created using either [connect()](#) or [dbConnect()](#).
sql	The SQL to be send.
errorReportFile	
	The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.
snakeCaseToCamelCase	
	If true, field names are assumed to use snake_case, and are converted to camelCase.
oracleTempSchema	
	DEPRECATED: use `tempEmulationSchema` instead.
tempEmulationSchema	
	Some database platforms like Oracle and Impala do not truly support temp tables. To emulate temp tables, provide a schema with write privileges where temp tables can be created.
integerAsNumeric	
	Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.
integer64AsNumeric	
	Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.
...	Parameters that will be used to render the SQL.

Details

This function calls the `render` and `translate` functions in the `SqlRender` package before calling [querySql()](#).

Value

A data frame.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
persons <- renderTranslatequerySql(conn,
  sql = "SELECT TOP 10 * FROM @schema.person",
  schema = "cdm_synpuf"
)
disconnect(conn)

## End(Not run)
```

---

renderTranslateQuerySqlToAndromeda

Render, translate, and query to local Andromeda

Description

This function renders, and translates SQL, sends it to the server, and returns the results as an ffdf object

Usage

```
renderTranslateQuerySqlToAndromeda(
  connection,
  sql,
  andromeda,
  andromedaTableName,
  errorReportFile = file.path(getwd(), "errorReportSql.txt"),
  snakeCaseToCamelCase = FALSE,
  appendToTable = FALSE,
  oracleTempSchema = NULL,
  tempEmulationSchema = getOption("sqlRenderTempEmulationSchema"),
  integerAsNumeric = getOption("databaseConnectorIntegerAsNumeric", default = TRUE),
  integer64AsNumeric = getOption("databaseConnectorInteger64AsNumeric", default = TRUE),
  ...
)
```

Arguments

connection	The connection to the database server created using either `connect()` or `dbConnect()`.
sql	The SQL to be send.
andromeda	An open Andromeda object, for example as created using `Andromeda::andromeda()`.
andromedaTableName	
	The name of the table in the local Andromeda object where the results of the query will be stored.

errorReportFile

> The file where an error report will be written if an error occurs. Defaults to 'errorReportSql.txt' in the current working directory.

snakeCaseToCamelCase

> If true, field names are assumed to use snake_case, and are converted to camel-Case.

appendToTable If FALSE, any existing table in the Andromeda with the same name will be replaced with the new data. If TRUE, data will be appended to an existing table, assuming it has the exact same structure.

oracleTempSchema

> DEPRECATED: use `tempEmulationSchema` instead.

tempEmulationSchema

> Some database platforms like Oracle and Impala do not truly support temp tables. To emulate temp tables, provide a schema with write privileges where temp tables can be created.

integerAsNumeric

> Logical: should 32-bit integers be converted to numeric (double) values? If FALSE 32-bit integers will be represented using R's native `Integer` class.

integer64AsNumeric

> Logical: should 64-bit integers be converted to numeric (double) values? If FALSE 64-bit integers will be represented using `bit64::integer64`.

... Parameters that will be used to render the SQL.

Details

This function calls the `render` and `translate` functions in the `SqlRender` package before calling querySqlToAndromeda().

Value

Invisibly returns the andromeda. The Andromeda object will have a table added with the query results.

Examples

```
## Not run:
connectionDetails <- createConnectionDetails(
  dbms = "postgresql",
  server = "localhost",
  user = "root",
  password = "blah",
  schema = "cdm_v4"
)
conn <- connect(connectionDetails)
renderTranslatequerySqlToAndromeda(conn,
  sql = "SELECT * FROM @schema.person",
  schema = "cdm_synpuf",
  andromeda = andromeda,
  andromedaTableName = "foo"
)
disconnect(conn)

andromeda$foo
```

```
## End(Not run)
```

---

requiresTempEmulation   *Does the DBMS require temp table emulation?*

Description

Does the DBMS require temp table emulation?

Usage

```
requiresTempEmulation(dbms)
```

Arguments

dbms   The type of DBMS running on the server. See `connect()` or `createConnectionDetails()`
                for valid values.

Value

TRUE if the DBMS requires temp table emulation, FALSE otherwise.

Examples

```
requiresTempEmulation("postgresql")
requiresTempEmulation("oracle")
```

---

show,DatabaseConnectorDriver-method
                        Show an Object

Description

Display the object, by printing, plotting or whatever suits its class. This function exists to be specialized by methods. The default method calls `showDefault`.

Formal methods for show will usually be invoked for automatic printing (see the details).

Usage

```
## S4 method for signature 'DatabaseConnectorDriver'
show(object)
```

Arguments

object   Any R object

Details

Objects from an S4 class (a class defined by a call to [setClass](#)) will be displayed automatically is if by a call to show. S4 objects that occur as attributes of S3 objects will also be displayed in this form; conversely, S3 objects encountered as slots in S4 objects will be printed using the S3 convention, as if by a call to [print](#).

Methods defined for show will only be inherited by simple inheritance, since otherwise the method would not receive the complete, original object, with misleading results. See the simpleInheritanceOnly argument to [setGeneric](#) and the discussion in [setIs](#) for the general concept.

Value

show returns an invisible NULL.

See Also

[showMethods](#) prints all the methods for one or more functions.

---

year   *Extract the year from a date*

Description

This function is provided primarily to be used together with dbplyr when querying a database. It will also work in dplyr against data frames.

Usage

```
year(date)
```

Arguments

date   The date.

Value

The year

Examples

```
year(as.Date("2000-02-01"))
```

Index